



# ByteZip: Efficient Lossless Compression for Structured Byte Streams Using DNNs

**Parvathy Ramakrishnan P and Satyajit Das**

**Department of Data Science, Indian Institute of Technology, Palakkad**

July 1, 2024

# Outline of the Presentation

## 1 Introduction

- Overview
- Deep Learning in Lossless Data Compression
- Related Works
- Limitations

## 2 ByteZip


- Proposed Method
- Architecture
- Experiments and Results

## 3 Conclusion

## 4 References


# Data Compression

- Science of representing information in a compact form.



# Data Compression


- Science of representing information in a compact form.



- Lossless and Lossy data compression

# Data Compression


- Science of representing information in a compact form.



- Lossless and Lossy data compression
- Compression Ratio(size reduction) vs Compression Speed(runtime)

# Data Compression


- Science of representing information in a compact form.



- Lossless and Lossy data compression
- Compression Ratio(size reduction) vs Compression Speed(runtime)
- Nature of the data being compressed, tolerance for loss, and requirements for speed and memory efficiency.

# Data Compression

- Science of representing information in a compact form.



- Lossless and Lossy data compression
- Compression Ratio(size reduction) vs Compression Speed(runtime)
- Nature of the data being compressed, tolerance for loss, and requirements for speed and memory efficiency.
- Byte streams - Binary representation of information transmitted or stored as a continuous stream of bytes.

# Data Compression

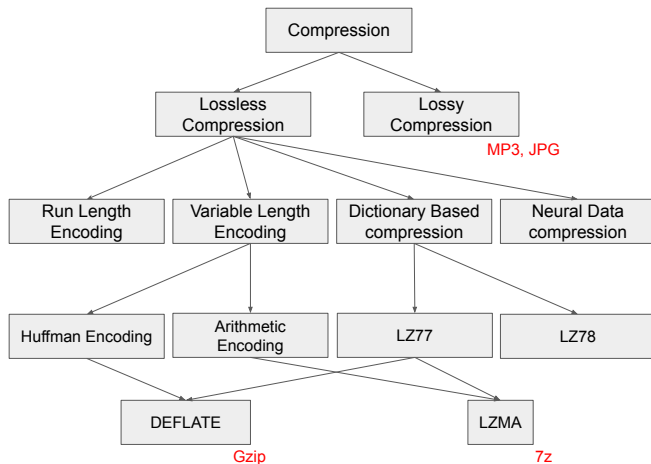


Figure: An overview of different data compression techniques



# Deep Learning in Lossless Data Compression

- Neural Data Compression<sup>1</sup> - Application of neural networks and other machine learning methods to data compression

---

<sup>1</sup>Yang, Mandt, and Theis, “An introduction to neural data compression”, 2022.

# Deep Learning in Lossless Data Compression

- Neural Data Compression<sup>1</sup> - Application of neural networks and other machine learning methods to data compression
- First build a probabilistic model of the data, and then feed its probabilities into an entropy coding scheme that converts data into compact bit-strings

---

<sup>1</sup>Yang, Mandt, and Theis, “An introduction to neural data compression”, 2022.

# Deep Learning in Lossless Data Compression

- Neural Data Compression<sup>1</sup> - Application of neural networks and other machine learning methods to data compression
- First build a probabilistic model of the data, and then feed its probabilities into an entropy coding scheme that converts data into compact bit-strings




Figure: Overview of Neural Data Compression

<sup>1</sup>Yang, Mandt, and Theis, “An introduction to neural data compression”, 2022.

# State of The Art Deep Learning Compressors

- CMIX<sup>2</sup>

- Bitwise predictions using an ensemble of independent models
- Boosting over multiple compressor experts leads to improved performance
- lstm-compress and tensorflow-compress

- NNCP<sup>3</sup>

- Byte-wise predictions
- Improved compression ratio using multi-layer LSTMs and Transformer XL models

---

<sup>2</sup>Knoll, *Cmix*, 2014; Knoll, *Lstm-compress: data compression using LSTM*, 2019; Knoll, *Tensorflow-compress: lossless data compression using neural networks in TensorFlow*, 2020.

<sup>3</sup>Bellard, “Lossless data compression with neural networks”, 2019; Bellard, “NNCP v2: Lossless Data Compression with Transformer”, 2021.

# State of The Art Deep Learning Compressors

- TRACE<sup>4</sup>
  - Addresses the execution time challenges associated with deep-learning-based compressors
  - Single layer transformer architecture with reduced number of parameters and better GPU utilization
  - 65% better compression ratio, but still is 1000x slower than Gzip

---

<sup>4</sup>Mao et al., “TRACE: A Fast Transformer-Based General-Purpose Lossless Compressor”, 2022.

# Limitations

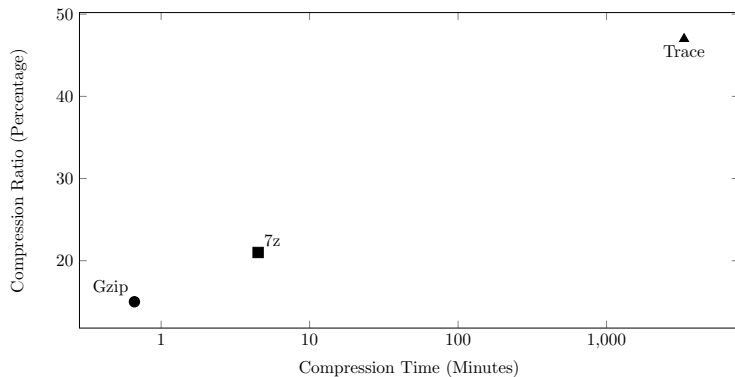


Figure: Compression Time vs Compression Ratio for compressing 1 GB data.

# Autoregressive Compressors

- Key Limitation- Very high runtime

# Autoregressive Compressors

- Key Limitation- Very high runtime
- 2 key factors
  - Complexity of the model used in probability modeling
  - Sequential nature of the overall compression process



# Autoregressive Compressors

- Key Limitation- Very high runtime
- 2 key factors
  - Complexity of the model used in probability modeling
  - Sequential nature of the overall compression process
- Sequential GPU/CPU/IO operations for each byte

# Autoregressive Compressors

- Key Limitation- Very high runtime
- 2 key factors
  - Complexity of the model used in probability modeling
  - Sequential nature of the overall compression process
- Sequential GPU/CPU/IO operations for each byte

Table: Average Compression and Decompression Speed

Model	Avg Comp Speed (MB/Min)	Avg Decomp Speed (MB/Min)
TRACE_8	2.01	1.39
FNN_8	2.92	1.73
Uniform AE	3.31	1.83

- 1 Introduction
- 2 ByteZip
  - Proposed Method
  - Architecture
  - Experiments and Results
- 3 Conclusion
- 4 References

# Overview

- Hierarchical autoencoder-based approach<sup>5</sup> to model probability for byte streams instead of existing autoregressive approaches.
- Instead of a forward pass per byte, a chunk of data is processed in one pass resulting in faster compression.
- Support for partial decompression

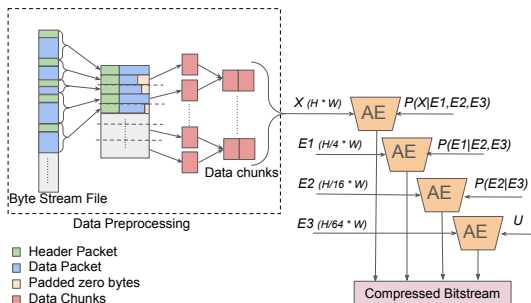



Figure: Overview of ByteZip compressor

<sup>5</sup>Mentzer et al., "Practical Full Resolution Learned Lossless Image Compression", 2019.

# Data Preprocessing

Transforms the input data into a form that is more easily compressible



Byte Stream File

## Data Preprocessing

- |               |                   |
|---------------|-------------------|
| Header Packet | Padded zero bytes |
| Data Packet   | Data Chunks       |

---

### Algorithm 1 Data Preprocessing

---

```
1: procedure PREPROCESSDATA(bytestream)
2:   for all packet in bytestream do
3:     if  $n$  repetitive packets then
4:       Join  $n$  packets to a single record
5:        $w \leftarrow w_1 + w_2 + \dots + w_n$ 
6:     end if
7:   end for
8:    $W \leftarrow$  Width of longest record
9:   for all record in bytestream do
10:     $record \leftarrow record + [(W - w) \text{ zero bytes}]$ 
11:   end for
12:    $N \leftarrow \text{len}(\text{bytestream}) / (H * W)$ 
13:    $chunks \leftarrow \text{bytestream.reshape}(N, H, W)$ 
14:   if  $W < 256$  then
15:      $k \leftarrow \text{math.ceil}(256/W)$ 
16:      $chunks \leftarrow \text{chunks.reshape}(N/k, H, k * W)$ 
17:   end if
18: end procedure
```

---

# Probability Modelling

- An autoencoder with a mixed-density network is employed to model probability

# Probability Modelling

- An autoencoder with a mixed-density network is employed to model probability
- Three key modules - Encoder, Decoder and Probability estimator

# Probability Modelling

- An autoencoder with a mixed-density network is employed to model probability
- Three key modules - Encoder, Decoder and Probability estimator
- Leverage the lower dimensional representation during decompression to calculate probability parameters



# Probability Modelling

- An autoencoder with a mixed-density network is employed to model probability
- Three key modules - Encoder, Decoder and Probability estimator
- Leverage the lower dimensional representation during decompression to calculate probability parameters
- Multiscale hierarchical Autoencoders- Reduce the size of latent variables, Capture complex dependencies

# Probability Modelling

- An autoencoder with a mixed-density network is employed to model probability
- Three key modules - Encoder, Decoder and Probability estimator
- Leverage the lower dimensional representation during decompression to calculate probability parameters
- Multiscale hierarchical Autoencoders- Reduce the size of latent variables, Capture complex dependencies
- Trained to minimize the sum of negative log-likelihood in all scales

# Neural Network Architecture

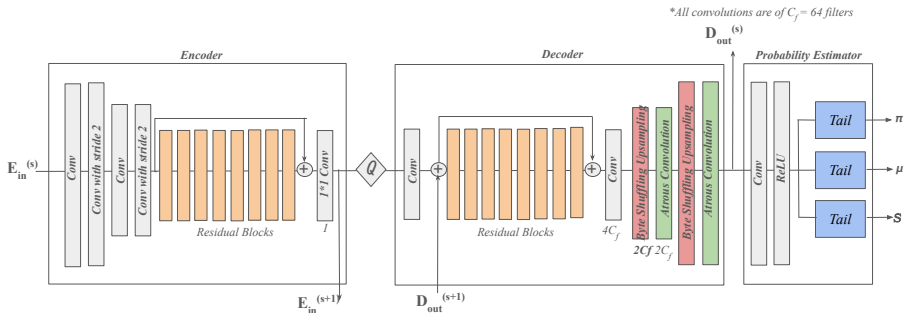


Figure: Architecture details for a single scale

# Neural Network Architecture

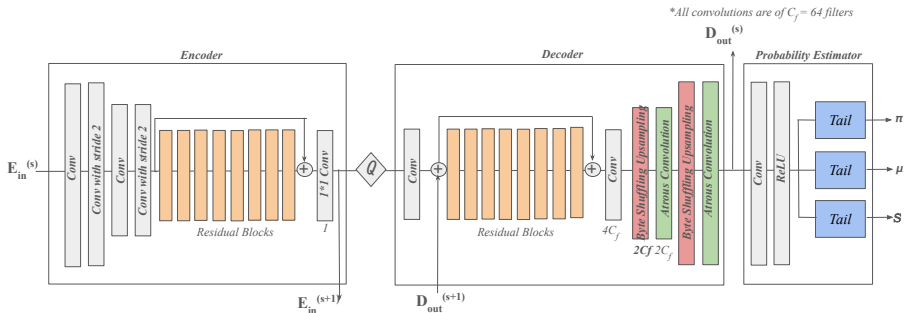


Figure: Architecture details for a single scale




Figure: Residual and Tail Blocks

# Neural Network Architecture

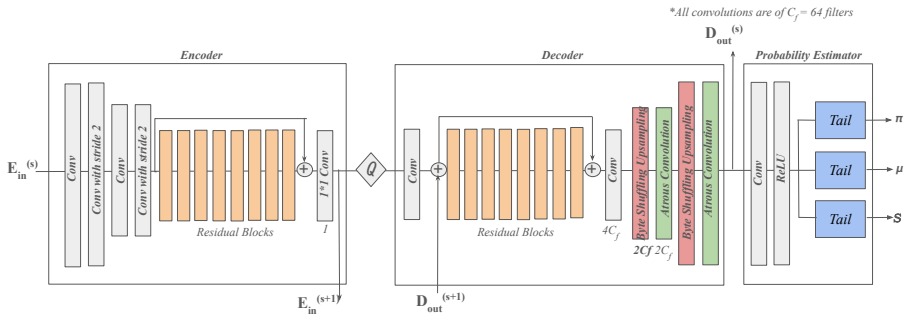


Figure: Architecture details for a single scale

$$P(x|\pi, \mu, s) \sim \sum_{i=1}^K \pi_i \text{logistic}(\mu_i, s_i)$$

# Neural Network Architecture


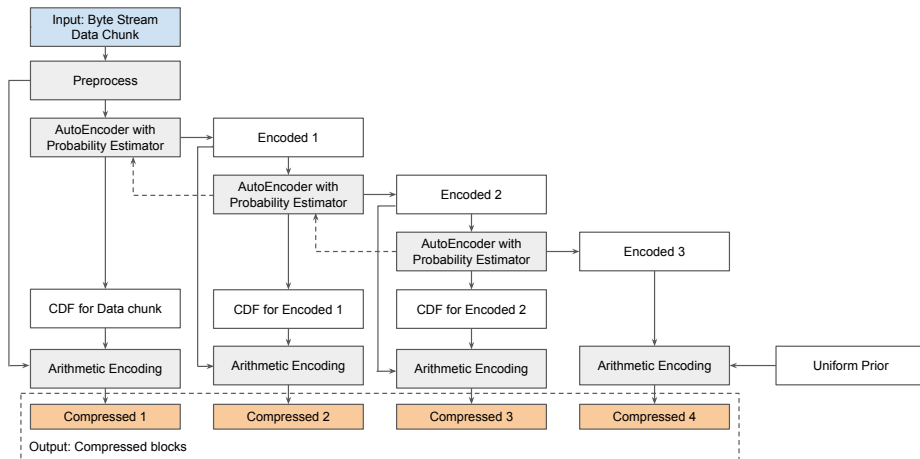


Figure: Architecture details for a single scale

$$P(x|\pi, \mu, s) \sim \sum_{i=1}^K \pi_i \text{logistic}(\mu_i, s_i)$$

$$P(x|\pi, \mu, s) \sim \sum_{i=1}^K \pi_i [\sigma((x + b/2 - \mu_i)/s_i) - \sigma((x - b/2 - \mu_i)/s_i)]$$

# Compression



### Figure: Hierarchical Compression Framework

# Decompression

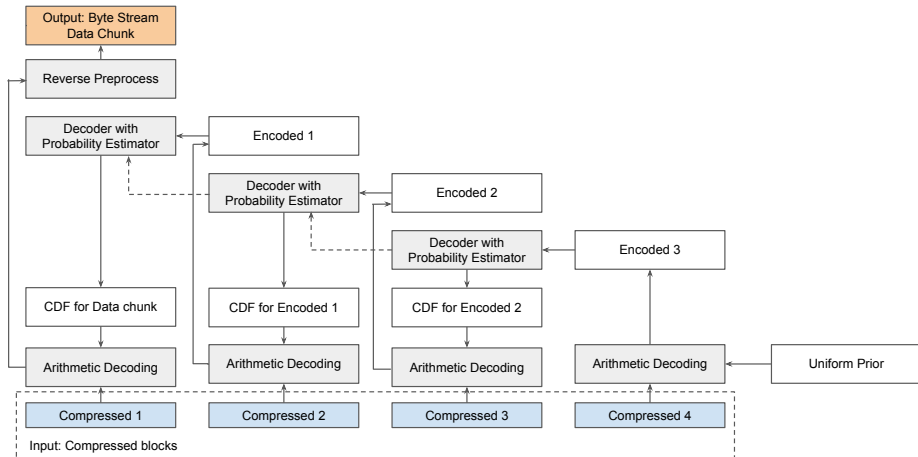


Figure: Hierarchical Decompression Framework



# Experimental Setup

## Models considered for performance comparison:

Model	Type	Pretraining Required	Parameter Updation During Compression	Support Partial Decompression	CPU/GPU
Gzip <sup>6</sup>	Traditional	No	No	No	CPU
7z <sup>7</sup>	Traditional	No	No	No	CPU
TRACE <sup>8</sup>	Autoregressive	No	Yes	No	GPU
FNN	Autoregressive	No	Yes	No	GPU
ByteZip	Non Autoregressive	Yes	No	Yes	GPU

---

<sup>6</sup>Deutsch, *GZIP file format specification version 4.3*, 1996.

<sup>7</sup>7-Zip.

<sup>8</sup>Mao et al., “TRACE: A Fast Transformer-Based General-Purpose Lossless Compressor”, 2022.

# Experimental Setup

## Models considered for performance comparison:

Model	Type	Pretraining Required	Parameter Updation During Compression	Support Partial Decompression	CPU/GPU
Gzip <sup>6</sup>	Traditional	No	No	No	CPU
7z <sup>7</sup>	Traditional	No	No	No	CPU
TRACE <sup>8</sup>	Autoregressive	No	Yes	No	GPU
FNN	Autoregressive	No	Yes	No	GPU
ByteZip	Non Autoregressive	Yes	No	Yes	GPU

## Evaluation metrics:

- Average Compression Speed =  $\frac{\text{Data Size (MB)}}{\text{Total Compression Time (Min)}}$
- Compression Ratio =  $\frac{\text{Original Size}}{\text{Compressed Size}}$
- Compression Percentage =  $\left(1 - \frac{\text{Compressed Size}}{\text{Original Size}}\right) \times 100$

---

<sup>6</sup>Deutsch, *GZIP file format specification version 4.3*, 1996.

<sup>7</sup>7-Zip.

<sup>8</sup>Mao et al., "TRACE: A Fast Transformer-Based General-Purpose Lossless Compressor", 2022.

# Experimental Setup

## Hardware:

Single workstation with Intel(R) Core(TM) i9-10900K CPU @ 3.70GHz, NVIDIA GeForce RTX3080 Ti GPU with 10240 cores and 12GB RAM

---

<sup>9</sup>Burtscher and Ratanaworabhan, “FPC: A High-Speed Compressor for Double-Precision Floating-Point Data”, 2009.

<sup>10</sup>Burtscher and Ratanaworabhan, “FPC: A High-Speed Compressor for Double-Precision Floating-Point Data”, 2009.

<sup>11</sup>Choi et al., “Children’s song dataset for singing voice research”, 2020.

# Experimental Setup

## Hardware:

Single workstation with Intel(R) Core(TM) i9-10900K CPU @ 3.70GHz, NVIDIA GeForce RTX3080 Ti GPU with 10240 cores and 12GB RAM

## Dataset:

Table: Description of datasets along with the parameters used

Name	Dataset Size	Description	chunk size	k	scale
Sonar	25 GB	Ocean acoustic data collected using sonar arrays	64 * 1504	3	3
obs_info <sup>9</sup>	500 MB	Latitude and Longitude information of the observation points of a weather satellite	256 * 256	3	4
obs_spitzer <sup>10</sup>	1.1 GB	Spitzer Space Telescope data with 64-Bit Double-Precision Floating-Point values	256 * 256	2	4
CSD <sup>11</sup>	1.5 GB	Children's Song Dataset for Singing Voice Research	256 * 256	3	4

---

<sup>9</sup>Burtscher and Ratanaworabhan, "FPC: A High-Speed Compressor for Double-Precision Floating-Point Data", 2009.

<sup>10</sup>Burtscher and Ratanaworabhan, "FPC: A High-Speed Compressor for Double-Precision Floating-Point Data", 2009.

<sup>11</sup>Choi et al., "Children's song dataset for singing voice research", 2020.

# Runtime Improvements




Figure: Compression Time




Figure: Decompression Time

# Runtime Improvements




Figure: Compression Time




Figure: Decompression Time

Table: Average Compression and Decompression Speed

	Model	Avg Comp Speed (MB/Min)	Avg Decomp Speed (MB/Min)
Traditional	Gzip	1756	10974
	7z	191	1895
Non-Autoregressive	ByteZip(Ours)	20	18
Autoregressive	FNN_64	2.89	1.76
	TRACE_8	2.01	1.34
	TRACE_64	0.31	0.28

# Size Reduction




Figure: Compression percentage on different filesizes

# Size Reduction




Figure: Compression percentage on different filesizes

Table: Average compression ratio

	Model	obs_info	obs_spitzer	CSD	Sonar
Traditional	Gzip	1.15	1.06	1.51	1.20
	7z	1.22	1.13	2.19	1.30
Non-Autoregressive	ByteZip(Ours)	2.17	1.21	3.01	1.65
Autoregressive	FNN_64	2.09	1.19	1.67	1.57
	TRACE_8	2.01	1.18	2.92	1.40
	TRACE_64	2.81	1.28	3.08	1.84



# Compression Time vs Compression Ratio

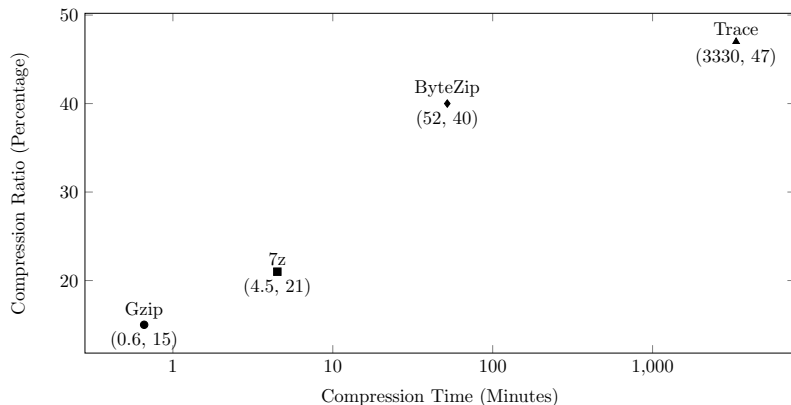


Figure: Compression Time vs Compression Ratio for compressing 1 GB bytestream data

1 Introduction

2 ByteZip









3 Conclusion

4 References

- Autoregressive bitwise prediction-based lossless compression models have limitations in improving runtime
- Our proposed method ByteZip balances the higher compression ratio achieved by autoregressive neural network models with the practicality of attaining a reasonable compression speed
- Future Scope:
  - Reduce the computational requirements of the neural network compressor to create a more practical compressor.
  - Support the incorporation of domain-specific knowledge into the compressor when available.

Thank You.

# References I

-  Yang, Yibo, Stephan Mandt, and Lucas Theis. “An introduction to neural data compression”. In: *arXiv preprint arXiv:2202.06533* (2022).
-  Knoll, Bryon. *Cmix*. <https://www.byronknoll.com/cmixon.html>. 2014.
-  —. *Lstm-compress: data compression using LSTM*. <https://github.com/byronknoll/lstm-compress>. 2019.
-  —. *Tensorflow-compress: lossless data compression using neural networks in TensorFlow*. <https://github.com/byronknoll/tensorflow-compress>. 2020.
-  Bellard, Fabrice. “Lossless data compression with neural networks”. In: *URL: https://bellard.org/nncp/nncp.pdf* (2019).
-  —. “NNCP v2: Lossless Data Compression with Transformer”. In: (2021).
-  Mao, Yu et al. “TRACE: A Fast Transformer-Based General-Purpose Lossless Compressor”. In: *Proceedings of the ACM Web Conference 2022*. ACM, 2022, 1829–1838. ISBN: 9781450390965.
-  Mentzer, F. et al. “Practical Full Resolution Learned Lossless Image Compression”. In: *CVPR*. 2019.

# References II



Deutsch, Peter. *GZIP file format specification version 4.3*. Tech. rep. 1996.



7-Zip.



Burtscher, Martin and Paruj Ratanaworabhan. “FPC: A High-Speed Compressor for Double-Precision Floating-Point Data”. In: *IEEE Transactions on Computers* 58.1 (2009), pp. 18–31. DOI: 10.1109/TC.2008.131.



Choi, Soonbeom et al. “Children’s song dataset for singing voice research”. In: *International Society for Music Information Retrieval Conference (ISMIR)*. 2020.